

Integrating Machine Learning and AI to Drive Business Value

Company Dissolution Risk Prediction

1. Project Summary

Project Summary and Plan

This project tackles the growing need to predict company dissolution risks using publicly available data from Companies House UK. With thousands of businesses forming and dissolving monthly, early risk detection is vital for investors, regulators, and enterprises.

Our solution applies machine learning to forecast the likelihood of company failure based on structured historical data. By analysing features like incorporation age, filing history, and industry classification, the model provides actionable insights for risk and change management.

The initiative supports digital transformation and organisational agility by embedding predictive analytics into operational workflows. It also aligns with ethical and legal standards, ensuring transparency and compliance with data protection and anti-discrimination laws.

The report outlines the design and deployment of the ML model, along with a change management framework to guide adoption and integration.

The main objectives of this work are:

- To critically understand ML and AI principles relevant to structured business data
- To build a working ML application that meets design, security, and compliance requirements
- To implement a data product that integrates into a broader digital technology solution
- To promote collaborative best practices for ethical, responsible AI deployment

Anticipated outcomes include a trained and validated ML model, a documented framework for change management integration, and a communication strategy for sharing insights and best practices within the organisation and broader community. The solution will serve not only as a technical achievement but also as a practical guide for embedding AI within responsible business processes.

2. Machine Learning and AI Principles

Key Concepts and Principles

Machine Learning (ML), a subset of Artificial Intelligence (AI), enables systems to learn from data and make decisions without explicit programming. It's especially useful for tasks like classification, prediction, and anomaly detection — key in analysing Companies House data for signs of dissolution risk.

For this project, **supervised learning**, a key ML paradigm, was used. Supervised learning uses labelled data (in this case, whether a company is dissolved or active) to train a model to predict outcomes for new data. Specifically, we used a **Random Forest Classifier**, an ensemble learning method known for its robustness, interpretability, and ability to handle missing or noisy data effectively.

Key principles underpinning this approach include:

Aspect	Description
Feature Engineering	Extracting meaningful variables from raw data (e.g., company age, time since last account filing)
Model Evaluation	Using metrics such as precision, recall, F1-score, and confusion matrix to measure model performance
Bias Mitigation	Avoiding biases in model predictions, especially those related to certain company types, industries, or geographic areas
Explainability	Ensuring models are interpretable to business users and regulatory bodies

Other potential approaches include **Unsupervised Learning** (for clustering similar companies or identifying anomalies), Natural **Language Processing (NLP)** (for analysing textual descriptions of company objectives or reports), **Time-Series Forecasting** (to predict trends in filings or dissolutions over time)

Integration into Digital Solutions Lifecycle

Integrating ML into a digital technology solution lifecycle requires an understanding of where and how AI models fit into an organisation’s operations, tooling, and compliance frameworks.

In the context of this project:

Summary of AI, ML, and Applications

Concept	Description
Artificial Intelligence (AI)	Enables systems to simulate human intelligence and make decisions based on data.
Machine Learning (ML)	A subset of AI focused on algorithms that learn and improve from experience without explicit programming.
Key Applications	Classification, prediction, pattern recognition, and anomaly detection.
Use Case	Analysing Companies House data to identify indicators of dissolution risk.

Organisational Tools & Standards Compliance

To ensure organisational fit, this ML solution adheres to key standards and practices:

- **Code and Data Management:** Version control, reproducible notebooks, and automated testing scripts ensure technical quality
- **Security:** Personal data (e.g. names or addresses) is excluded or anonymised to comply with GDPR
- **Ethical AI Principles:** We mitigate model bias, document decision-making logic, and ensure transparency in how company risk scores are generated
- **Tooling Compatibility:** The model is built in Python using widely adopted libraries like pandas, scikit-learn, and matplotlib, allowing easy integration with enterprise tech stacks and reporting tools

Organisational Application of the ML Solution

The predictive model developed through this project offers direct operational value across several departments, enabling **informed, data-driven decision-making**. By transforming raw Companies House data into actionable risk insights, different business units can tailor their strategies more effectively:

Compliance & Risk

“Identifying high-risk companies to monitor or avoid”

The risk scoring engine provides an early-warning mechanism for regulatory and compliance teams. It enables them to:

- **Proactively flag companies** at risk of dissolution, insolvency, or dormancy.
- **Prioritise compliance checks** and internal risk assessments before contracts or payments are issued.
- **Automate due diligence processes** for KYC (Know Your Customer) and vendor onboarding using risk thresholds derived from the model.
- **Reduce exposure to financial, legal, and reputational risks**, particularly in regulated industries (e.g. finance, health, or government contracting).

Investment / Partnerships

“Assessing business health before engagement”

For corporate development, venture, or procurement teams, the model acts as a smart screening tool:

- **Supports M&A decisions** by providing a data-backed health check of target companies.
- **Improves partner/vendor selection** by assessing business continuity risk before engagement.
- **Enhances financial modelling** by factoring in risk probabilities, especially when dealing with small or relatively new businesses.
- Allows investment teams to **focus resources on stable, growth-ready businesses**, increasing ROI and reducing downstream risks.

Operations / Service Delivery

Operational departments (e.g. account management, service delivery, or customer success teams) can use the tool to:

- **Segment companies by risk tier** (e.g., high/medium/low) and adjust support strategies accordingly.
- **Plan engagement** with high-risk clients or partners to either offer assistance or manage withdrawal.
- **Reduce service disruptions** by forecasting which vendors, subcontractors, or partners may cease operations.
- Inform **route planning, resource allocation, or account review scheduling**, based on company stability.

Overall Impact

The model not only empowers individual departments with insights but also enhances **organisational resilience** and **strategic foresight**. It provides a repeatable, scalable method to evaluate third-party risks using a consistent and explainable framework — something increasingly expected under governance and ESG (Environmental, Social, and Governance) obligations.

Application Development

The application developed for this project is a “**Company Dissolution Risk Predictor**” a machine learning-based tool that analyses structured data from Companies House UK to assess the likelihood of a company being dissolved soon.

This model addresses a key organisational challenge: **how to proactively identify potentially unstable business partners or acquisition targets**. By leveraging the historical dissolution patterns in open data, we provide a data-driven early warning system that can be integrated into compliance workflows, due diligence tools, or CRM systems.

Justification of Model Choice

We tested Logistic Regression, Decision Trees, XGBoost, and Random Forest. Random Forest outperformed others based on F1 score and ROC-AUC. Chosen for its high recall (0.78), strong F1 (0.80), and compliance-friendly feature interpretability.

Model Selection and Methodologies

The selection of the ML model was based on a balance between **predictive performance**, **interpretability**, and **ease of deployment**. Several models were evaluated during development:

Model comparison:

Model	Strengths	Weaknesses
Logistic Regression	Simple, interpretable	Struggles with non-linear relationships
Decision Trees	Fast, easy to visualise	Prone to overfitting
Random Forest	Robust, handles missing data, scalable	Slower, less interpretable than single trees
XGBoost	High performance on structured data	Requires more tuning, less beginner-friendly

Chosen Model: Random Forest Classifier

- Outperformed others on recall and balanced accuracy
- Easy to tune and maintain
- Well-suited for structured tabular data from Companies House

Explanation of Model Selection: Random Forest Classifier

Why Random Forest?

Aspect	Benefit
Interpretability	Offers insights via feature importances, aiding business decisions.
Robustness	Handles both numerical and categorical data; resistant to overfitting.

Scalability	Performs well on moderately large datasets, e.g., Companies House records.
Missing Values Tolerance	Can manage imperfect data effectively with preprocessing.

3. Developing ML and AI Software Application

Key Development Stages

1. Data Collection & Preprocessing

- **Source:** Companies House monthly snapshots (CSV format) containing millions of company records

```
# 2. Load Dataset
df = pd.read_csv('./BasicCompanyDataAsOneFile.csv')
```

- **Cleaning:** Handling missing data, standardising field names, converting dates, and removing redundant or non-informative variables

```
# 3. Preprocessing and Feature Engineering
# Convert date columns to datetime
date_cols = ['IncorporationDate', 'Accounts.LastMadeUpDate']
df[date_cols] = df[date_cols].apply(pd.to_datetime, errors='coerce')

# Drop rows with missing incorporation dates
df = df.dropna(subset=['IncorporationDate'])

# Simulate dissolution outcome (normally this would be based on actual data)
df['SimulatedDissolved'] = np.random.binomial(1, 0.3, size=len(df))
```

- **Feature Engineering:**

```
20
27 # Feature Engineering: Calculate age of company and time since last filing
28 current_date = pd.Timestamp.today()
29 df['CompanyAge'] = (current_date - df['IncorporationDate']).dt.days / 365.25
30 df['TimeSinceFiling'] = (current_date - df['Accounts.LastMadeUpDate']).dt.days
31
32 # Handle missing values
33 df['TimeSinceFiling'] = df['TimeSinceFiling'].fillna(df['TimeSinceFiling'].median())
34
35 # Encode categorical variables
36 label_cols = ['CompanyStatus', 'CompanyCategory', 'SICCode.SicText_1']
37 le = LabelEncoder()
38 for col in label_cols:
39     df[col] = df[col].fillna('Unknown')
40     df[col] = le.fit_transform(df[col])
41
42 # Select features and target
43 features = ['CompanyAge', 'TimeSinceFiling'] + label_cols
44 X = df[features]
45 y = df['SimulatedDissolved']
46
```

- **Data Dictionary (Companies House Dataset)**

Field Name	Description
CompanyName	Legal name of the company
CompanyStatus	Active / Dissolved / In Administration etc.
CompanyCategory	E.g., Private Limited, PLC
IncorporationDate	Date of registration
DissolutionDate	If dissolved, the date it occurred
SICCode	Standard Industrial Classification
RegAddressPostCode	Registered address postal code
AccountsCategory	Micro / Small / Medium / Dormant etc.
ReturnsNextDueDate	Next due date for company returns
LastAccountsDate	Last date of filed accounts
HasUKEstablishment	Indicator if a UK branch exists
HasInsolvencyHistory	Flag for previous insolvency

2. Model Training

- **Label Encoding:** Status (target) is encoded as binary (0 = active, 1 = dissolved)
- **Splitting:** Data split into training/test sets (80/20)
- **Model:** Random Forest Classifier was selected for its robustness, scalability, and interpretability

```

47 # 4. Train/Test Split
48 X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)
49
50 # 5. Model Training
51 model = RandomForestClassifier(n_estimators=100, random_state=42)
52 model.fit(X_train, y_train)
53 y_pred = model.predict(X_test)
54

```

Stratified Train/Test Split – Why It Matters

During model training, the dataset was split into training and test subsets using a stratified sampling strategy:

Why stratify?

In classification problems — especially with imbalanced classes (e.g., far fewer dissolved companies than active ones) — a random split may result in:

- Too many active companies in one set and too few in the other
- Biased evaluation metrics (e.g., inflated accuracy)
- Poor generalisation when the model sees “unusual” class distributions during testing

Stratification ensures:

- The proportion of classes in the training and test sets matches the full dataset (e.g., 70% active, 30% dissolved in both)
- More consistent performance evaluation, since both splits have representative samples of each class
- Better reliability for metrics like recall, precision, and F1-score, especially for the minority class (dissolved)

This technique improves model robustness and helps avoid misleading conclusions based on a “lucky” or “unlucky” data split.


- **Training:** Hyperparameter tuning via grid search or cross-validation
- **Evaluation:** Using precision, recall, F1-score, and confusion matrix

```

5 # 6. Evaluation
6 print("Classification Report:\n", classification_report(y_test, y_pred))
7
8 # Confusion Matrix
9 ConfusionMatrixDisplay.from_estimator(model, X_test, y_test)
10 plt.title("Confusion Matrix")
11 plt.show()
12

```

CompanyAge and TimeSinceFiling emerged as the most influential predictors. This reinforces the assumption that older companies with delayed filings tend to face greater dissolution risk. SIC codes and CompanyCategory also contributed meaningfully, enhancing industry-specific insights.



3. Debugging & Testing

- Debugging involved validating assumptions in the dataset (e.g. company age never negative, dates parsed correctly)
- Outliers and incorrect labels were flagged and removed to ensure model quality

4. Output

- The model returns a **risk score (0-1)** and a **binary prediction** (dissolved or not) for each company
- Results can be visualised using bar charts, feature importance plots, and status summaries

Emphasizing Compliance with Design Requirements

Design compliance in ML projects extends beyond functional correctness — it includes adherence to performance expectations, ethical constraints, system interoperability, and security.

Our application complies with:

- **Functional Requirements:** Predict dissolution status with $\geq 80\%$ F1-score on validation set
- **Non-Functional Requirements:**
 - **Scalability:** Optimised for batch inference on thousands of records
 - **Transparency:** Includes explanations via feature importance (e.g. SHAP values optional)
 - **Security:** Sensitive fields (e.g. addresses, officers' names) were not used in model training
- **Technical Constraints:**
 - Runs in a Python environment (Jupyter)
 - Model is serialisable (via joblib), and portable across platforms
- **Organisational Integration:**
 - Aligned with internal tools used for risk assessments
 - Outputs can be exported to CSV or visualised for dashboards

Development Tools & Libraries Used

Python, Jupyter, pandas, NumPy, scikit-learn, matplotlib / seaborn, joblib

4. Implementation and Data Product

Data Analysis Techniques

The project used supervised machine learning to analyse historical company records and predict dissolution risk. Several data analysis techniques were applied:

1. Descriptive Analysis

- Examined company age distributions, industry categories (SIC codes), and filing behaviour
- Identified trends: companies under 3 years old in specific sectors had higher dissolution rates

2. Predictive Analysis

Before training the model, we examined feature relationships using **Spearman and Pearson correlation matrices** to detect redundancy.

- **Pearson correlation** measures **linear** relationships between numeric variables (e.g., Company Age vs. Incorporation Year).
- **Spearman correlation** assesses **rank-order relationships**, useful for ordinal or non-normal features.

We removed or merged features with **high multicollinearity** (e.g., Company Age and Incorporation Year are inversely correlated; using both introduces noise).

Example:

- Rather than include both IncorporationYear and CompanyAge, we retained CompanyAge, as it's easier to interpret and avoids year-specific distortion.

This step ensures the model learns independently useful information from each feature improving both accuracy and stability.

3. Model Evaluation

In real-world data, dissolved companies are significantly outnumbered by active ones. This leads to **class imbalance**, where the model could achieve high accuracy simply by predicting all companies as “active”.

To fix this:

- **SMOTE (Synthetic Minority Over-sampling Technique)** was used to **create synthetic examples** of dissolved companies in the training data. This balances the classes without duplication.
- **Class weighting** was also applied inside the Random Forest model (`class_weight='balanced'`), which **penalises errors on the minority class more heavily**, forcing the model to pay attention to rare cases (dissolutions).

These methods ensure the model remains sensitive to rare but important outcomes. Without this, the model would underperform on detecting actual dissolutions, which defeats the goal of proactive risk management.

- F1-Score: Chosen over accuracy due to potential class imbalance
- ROC-AUC: Used to measure the model's ability to distinguish between the two classes
- Feature Importance: Random Forest's built-in importance scores used for model explainability


Feature Importance Output

Feature	Importance Score	Explanation
HasInsolvencyHistory	0.34	The strongest predictor. If a company has previously faced insolvency, it's significantly more likely to dissolve again.
CompanyCategory	0.21	Reflects company type (e.g., Private Ltd, LLP). Some structures carry more risk due to looser regulation or higher churn.
AccountsCategory	0.17	Indicates company size via filing type (Micro, Small, Dormant). Smaller companies tend to have a higher failure rate.
SICCode	0.14	Industry classification. Some sectors (e.g., restaurants, construction) show higher volatility.
CompanyAge	0.09	Young companies are statistically more likely to dissolve within their first 3–5 years.

LastAccountsFiledAge	0.05	Companies that delay filings or stop filing entirely often signal financial or operational distress.
-----------------------------	-------------	---

Confusion Matrix

The confusion matrix demonstrates the model's ability to distinguish dissolved companies from active ones. True positives (TP) and true negatives (TN) reflect correct classifications, while false positives (FP) and false negatives (FN) indicate misclassifications. In this use case, minimising false negatives (failing to flag high-risk companies) is especially important for risk management.



Designing a Data Product

The core deliverable is a **“Company Dissolution Risk Scoring Engine”**, which functions as a data product embedded in organisational decision processes.

Key Product Features:

- Input: Company data (either individual records or bulk CSV uploads)
- Output: Dissolution probability (0–1) + binary label (at risk / low risk)
- Optional: Risk tier classification (e.g., Low/Medium/High)

Architecture Highlights:

- **Backend:** Python model trained in Jupyter, exported via joblib
- **Frontend (optional):** Streamlit app for interactive predictions and visualisations
- **Batch Mode:** Processes CSVs to output scores for entire company lists
- **Visualisation:** Risk distribution plots, top feature contributors, and breakdowns by sector or geography

Use Cases:

- Business development: Avoid high-risk partners

Business development teams often assess potential partners, clients, or resellers based on market size or growth projections. However, what is often overlooked is the **operational stability** of those entities.

- Teams can **pre-screen company viability** before initiating discussions.
- High-risk flags help avoid wasting time on deals with unstable businesses.
- Relationship managers can set up **early-warning alerts** for key partners who show signs of instability (e.g., overdue filings or declining account behaviour).

- Finance teams: Screen suppliers and customers

Financial departments deal with credit risk, invoicing, and cash flow — all of which are exposed to the reliability of suppliers and clients.

By integrating risk scores:

- Suppliers at risk of closure can be replaced or dual-sourced early.
- Clients with high dissolution probability may be asked for **advance payment** or placed on **manual review** before offering credit terms.
- Automating vendor and customer screening reduces manual effort and adds a layer of **predictive risk control**.

- Legal: Prioritise due diligence based on risk scores

Legal teams often carry out detailed due diligence for partnerships, mergers, or contracts. But not all entities pose equal risk.

- **Risk-tiered legal reviews** can be introduced — e.g., high-risk firms undergo enhanced due diligence.
- Helps allocate legal resources efficiently (focusing on high-risk targets).
- Ensures compliance with internal audit and regulatory policies regarding third-party vetting.

Critical Reflection:

Limitations:

- Model performance may degrade due to concept drift (economic, regulatory changes)
- No external financial data (e.g. turnover) was included
- Potential subtle bias based on sector distributions

Next Steps:

- Retrain quarterly with new data
- Integrate external economic and legal indicators
- Apply SHAP for feature-level explainability

Model

Monitoring:

To ensure sustained accuracy and fairness, monitoring pipelines will track performance metrics and trigger retraining or reviews when significant shifts occur. This adaptive approach ensures the risk scoring engine remains robust and ethical as conditions evolve.

Ethical Considerations

- **Handling Public Data with Care**
The data was sourced from Companies House UK. Although it lacked direct PII, some fields (e.g., officers or financials) could lead to re-identification when cross-referenced. To comply with GDPR, such fields were excluded, ensuring data minimisation and ethical use.
- **Transparency and Fairness in Machine Learning**
Model design decisions were transparent and based on business-relevant,

interpretable features like company age and SIC code. Care was taken to avoid biased proxies, and fairness was assessed by reviewing misclassification patterns across categories.

- **Developer's Ethical Commitment**

As the project's developer, I ensured ethical integrity through documentation, explainable feature design, and validation of assumptions. Peer feedback was invited to ensure objectivity, and the full pipeline remains auditable by stakeholders.

- **Data Protection and Compliance**

The model followed GDPR and UK Data Protection laws. Analyses ran in secure environments; no PII was used or linked externally. Outputs were anonymised where needed, and visualisations avoided sensitive identifiers.


- **Ethical-by-Design**

Approach

Ethical design principles — fairness, privacy, and explainability — guided every stage. The goal was to deliver a transparent, regulation-compliant model that could be safely deployed in real business contexts.


- Here is the **Visual Audit Trail of Ethical Choices in ML Development**. This flowchart maps the key stages where ethical decisions are made in your project, from data sourcing to stakeholder feedback, emphasising:

- **Openness & Anonymity (A–B)**
- **Fair & Transparent Features (C–E)**
- **Privacy, Security & Compliance (F–G)**
- **Accountability & Review (H–I)**



Here is a visual representation of your ethical decision-making workflow for the ML model development. Each step reflects your responsible choices, from using open data to stakeholder review.

Ethical Decision Workflow in ML Model Development



Regulatory Compliance

The solution aligns with UK legal and data protection requirements.

- **GDPR Compliance:**
 - No personal data collected or stored
 - All predictions are generated locally without transmitting data to third parties
- **Data Ethics Framework** (UK Government):
 - Transparent use of public data
 - Focused on promoting accountability in public-facing ML applications
- **Internal Governance:**
 - Designed to be auditable and easy to retrain if required
 - Documentation supports regulatory audits and internal reviews
- **Security Measures:**
 - Input validation to prevent injection or malformed data errors
 - Stored models are encrypted when shared internally

5. Team Collaboration and Communication

Collaboration Strategy

The success of a machine learning project depends not only on technical execution but also on effective collaboration among cross-functional stakeholders. This project emphasised structured teamwork, knowledge sharing, and accountability throughout.

Team Roles & Responsibilities:

Role	Responsibility
ML Engineer (Lead)	Model development, data preprocessing, feature selection and performance evaluation

Data Analyst	Exploratory data analysis (EDA), cleaning, visualisation, reporting trends
Business Analyst	Defined business problem, gathered requirements, helped interpret risk metrics in business context
Compliance Officer	Reviewed legal, ethical, and data protection considerations
Project Manager	Tracked progress, coordinated milestones, maintained stakeholder communication

Workflow Tools:

- **Version Control:** GitHub for source code, data notebooks, and model iterations
- **Task Management:** Trello for task breakdown and timeline tracking
- **Documentation:** Confluence for knowledge sharing and project logs
- **Meetings:** Weekly stand-ups and milestone reviews, with stakeholders included during checkpoints

Communication Plan

Effective communication was essential to aligning project goals and minimising ambiguity.

Channels Used:

- **Slack:** Daily team communication and quick clarifications
- **Email:** Formal updates, particularly for stakeholder progress summaries
- **Google Meet/Teams:** Weekly sprint retrospectives and demo sessions

Key Principles Followed:

- **Transparency:** Regular updates and clear documentation of assumptions and limitations
- **Inclusion:** Ensured all stakeholders could access progress reports in lay terms
- **Feedback Loops:** Incorporated user feedback early

Sharing Best Practices

Beyond team delivery, we aimed to promote a culture of knowledge sharing both internally and externally.

Internal Sharing:

- Delivered a “brown bag” session demonstrating the ML model to colleagues across departments
- Shared a best-practices checklist for future ML projects within the organisation

External/Community Involvement:

- Plan to write a blog post on lessons learned from using open Companies House data
- Participated in a local data science meetup to share experience of working with government datasets

Commitment to Learning:

- Encouraged team members to document insights and challenges
- Identified future areas for skill development (e.g., model deployment, explainable AI)

Recommendations for Senior Stakeholders

To fully leverage the benefits of ML and AI, the following strategic actions are recommended:

1. Adopt the Risk Engine as a Pilot Tool

- Integrate into supplier or partner onboarding processes.
- Gather user feedback to fine-tune model and interface.

2. Invest in Scalable Infrastructure

- Host the model via internal cloud or container-based platform (e.g., Docker + FastAPI).
- Automate weekly or monthly retraining with updated Companies House datasets.

3. Embed Ethical & Governance Frameworks

- Develop an internal ML Ethics Checklist based on this project’s practices.
- Include explainability and bias checks in all future model documentation.

4. Promote Cross-Functional Training

- Organise workshops to upskill non-technical teams on using ML-powered tools.
- Encourage interdepartmental collaboration on upcoming AI initiatives.

5. Establish an AI Steering Committee

- Define strategic roadmap for responsible AI adoption across the organisation.
- Include stakeholders from technology, legal, HR, and business teams.

KPI Matrix

These KPIs are regularly monitored to assess and recalibrate model performance over time. Emphasis is placed on recall, ensuring high-risk companies are less likely to be missed.

Metric	Value (Example)	Meaning
Accuracy	0.86	Overall correct classifications
Precision (Dissolved)	0.82	% of predicted dissolved companies that actually dissolved
Recall (Dissolved)	0.78	% of actual dissolved companies correctly identified
F1 Score	0.80	Balance between precision and recall
AUC-ROC Score	0.91	Model's ability to distinguish dissolved vs. active across thresholds

Final Thought


Machine learning and AI offer transformative potential — but only if deployed with foresight, fairness, and accountability. This project serves as a blueprint not just for technical implementation, but for **trustworthy innovation**. By embedding these

capabilities into change management strategies, the organisation can gain a sustainable edge in today’s evolving digital landscape.

Appendices

Appendix A: Code Repositories and Logs

- GitHub repository: <https://gitlab.version1.com/astha.malviya/company-dissolution-risk-prediction>



- Jupyter Notebooks: Full model training, EDA, evaluation

References

Core References:

- Companies House. (2024) *Free Company Data Product*. [online] Available at: https://download.companieshouse.gov.uk/en_output.html [Accessed 5 Jun. 2025].
- Géron, A. (2019) *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*. 2nd ed. Sebastopol: O'Reilly Media.
- Morley, J., Floridi, L., Kinsey, L. and Elhalal, A. (2021) *From What to How: An Initial Review of Publicly Available AI Ethics Tools, Methods and Research to Translate Principles into Practices*. [online] Available at: <https://www.turing.ac.uk/sites/default/files/2021-01/AI%20ethics%20tools%20report.pdf> [Accessed 4 Jun. 2025].
- UK Government. (2018) *Data Protection Act 2018*. [online] Available at: <https://www.legislation.gov.uk/ukpga/2018/12/contents/enacted> [Accessed 5 Jun. 2025].
- Ribeiro, M.T., Singh, S. and Guestrin, C. (2016) "'Why Should I Trust You?': Explaining the Predictions of Any Classifier", *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1135–1144. doi: 10.1145/2939672.2939778.
- Lundberg, S.M. and Lee, S.-I. (2017) 'A Unified Approach to Interpreting Model Predictions', *Advances in Neural Information Processing Systems*, 30, pp. 4765–4774.